

TERRAFORM:

Terraform is an IAC tool.

It helps in provisioning the servers. (It also provides some basic configuration features for servers.)

Best practice is to use terraform only for provisioning servers.

Installation of Software and configuration is left to Software configuration management tools like Chef, Ansible and Puppet.

Installing Terraform:

1. Download the zip archive in <https://developer.hashicorp.com/terraform/tutorials/aws-get-started/install-cli>
2. Unzip the archive.
3. Move the file **mv ~/Downloads/terraform /usr/local/bin/**

Creating resources in AWS Using Terraform:

1. Creating a sample terraform file.
2. All terraform files end with **tf extension (.tf)**

main.tf

Here is where we write all the configurations.

Provider is responsible for understanding API interactions and exposing resources.

To talk to any of the cloud form - provider tag is used.

#We must install the provider individually with respect to cloud provider we are using.

```
provider "aws" {  
    region = "us-east-1"  
}
```

CMD Command:

1. `cd /Users/vinoth/terraform/01-terraform-basics#initializing backend`
2. `terraform init`

Creating an IAM user in AWS using root credentials:

SERVICE -> IAM -> Access Management -> Users -> Create User

UserName : command_lineXXXXXXXX_user

AccessKey: AKIAT7XsXXHxIG

SecretAccessKey: /tP/XV+csnto/DUXiwwP+glmJm

Trying to set Environment variables in CMD: **export**

AWS_ACCESS_KEY_ID=AKIAT7XAPXXHxIG

export AWS_SECRET_ACCESS_KEY=tP/XV+csnto/DUXiwwP+glmJm

Creating AWS S3 Bucket using Terraform:

#creating manually in aws console

Services -> s3 -> Create Bucket -> BucketName(Unique across globe)

S3 - Simple Storage Cloud (Used to store files, backup.)

```
provider "aws" {  
    region = "us-east-1"  
}
```

#plan and execute

Whenever we need to create something we have to include in resource tag.

#resource = providernameresouretype and internal terraform name

```
resource "aws_s3_bucket" "my_s3_bucket" {
```

#name of the bucket in aws cloud

```
Bucket = "my-s3-bucket-vinothkumar-6"
```

```
}
```

In terraform two step execution approach is used.

1. **Plan**
2. **Execute**

Command :

#- will give the overall plan details.

terraform plan

#- execute command

terraform apply

#output

Apply complete! Resources: 1 added, 0 changed, 0 destroyed.

S3 bucket will be created.

After successful bucket creation **terraform.tfstate** file will be generated it will have json objects.

1. Even if we execute the **terraform apply** without doing any changes to the file, It will just **refresh the state** and won't perform any actions.

#output

Apply complete! Resources: 0 added, 0 changed, 0 destroyed.

Terraform recognizes nothing as changed with help of In **main.tf** terraform scripts, we specify the desired state.(we need a s3 bucket, 5 VM's)

Terraform States

#desired -KNOWN - ACTUAL

1. Terraform states are needed to track the dependencies between different resources.
2. Terraform.tfstate acts as cache to retrieve the information soon.

1. **Desired State** - Specifying what needs to be created like a S3 bucket, 5 Virtual Machines, etc.. in the cloud

2. **Known State** - Is the result of previous execution. (Known status of previous execution and resource creations. Like the file **terraform.tfstate**)
3. **Actual State** - State of Bucket in AWS.(Changes made on the bucket)

Whenever we do **terraform apply** it looks for information in **terraform.tfstate**

Whatever resources are configured, it will verify in the AWS if any information is changed in AWS. It is called refreshing state...

And it realizes that the desired state is same as Actual state.

#changing the main.tf file renaming the bucket from my-s3-bucket-vinothkumar-6 to my-s3-bucket-vinothkumar-7.

2. If we execute the **terraform apply** with any changes to the file, It will just **refresh the state** and know the changes.

If we are creating to **rename a bucket**, it will ask below for permission to **delete existing bucket and create a new one.**

#output

Plan: 1 to add, 0 to change, 1 to destroy.

Do you want to perform these actions? Yes

Apply complete! Resources: 1 added, 0 changed, 1 destroyed.

#changing the main.tf file. Enabling the versioning for of the bucket.

3. If we execute the main.tf file using terraform apply after including the versioning of the bucket.

```
resource "aws_s3_bucket" "my_s3_bucket"{
  bucket ="my-s3-bucket-vinothkumar-7"
  versioning {
    enabled = true
  }
}
```

Now, desired state = versioning of bucket, Known state is terraform.tfstate of previous execution, and the actual state is whatever present in AWS.

terraform apply

```
#output
versioning {
    ~ enabled    = false -> true
    # (1 unchanged attribute hidden)
}

# (2 unchanged blocks hidden)
}
```

Plan: 0 to add, 1 to change, 0 to destroy.

Do you want to perform these actions? Yes

Apply complete! Resources: 0 added, 1 changed, 0 destroyed.

Terraform is Declarative.

#query about current state

terraform console ->

#providertype_resource.internalTerraformName

> aws_s3_bucket.my_s3_bucket

All Bucket information can be found.

#to check if bucket versioning is enabled or out.

>aws_s3_bucket.my_s3_bucket.versioning[0].enabled

true

We can include the property in **main.tf**

```
output "my_s3_bucket_complete_details" {
    Value = aws_s3_bucket.my_s3_bucket.versioning[0].enabled
}
```

Command to execute output and see:

#refresh=false is enabled to avoid the comparison between actual state(s3 in AWS)

terraform apply -refresh=false

All bucket related information will be showed.

Creating AWS IAM User using Terraform:

**#Creating an IAM user in AWS console using root credentials:
SERVICE -> IAM -> Access Management -> Users -> Create User**

Changes in main.tf file to create a IAM user using terraform:

```
resource "aws_iam_user" "my_iam_user" {  
  name = "my_iam_user_vino"  
}
```

**#Command to save the terraform plan to a file
terraform plan -out iam.tfplan**

**#trying to execute the plan saved in cmd to create the user.
terraform apply iam.tfplan**

**#Output:
aws_iam_user.my_iam_user: Creating...
aws_iam_user.my_iam_user: Creation complete after 2s
[id=my_iam_user_vino]**

Apply complete! Resources: 1 added, 0 changed, 0 destroyed.

```
Outputting the i'am details:output "my_iam_user_details" {  
  value = aws_iam_user.my_iam_user  
}
```

**#command to check output of ram user details
terraform apply refresh=false**

Outputs:

```
my_iam_user_details = {  
  "arn" =  
  "arn:aws:iam::272763436187:user/my_iam_user_vino"  
  "force_destroy" = false  
  "id" = "my_iam_user_vino"  
  "name" = "my_iam_user_vino"  
  "path" = "/"  
  "permissions_boundary" = ""  
  "tags" = tomap(null) /* of string */  
  "tags_all" = tomap({})  
  "unique_id" = "AIDAT7APXXSNQOFJUU3J4"  
}
```

The **arn** (amazon resource notation) is the unique identification of any AWS resources.

Updating IAM username in terraform main.tf

We can directly change the name in resource it will delete the user and create a new user with new name.

```
resource "aws_iam_user" "my_iam_user" {  
  name = "my_iam_user_vinoth"  
}
```

#command

terraform plan -out iam.tfplan

Plan: 0 to add, 1 to change, 0 to destroy.

Changes to Outputs:

```
~ my_iam_user_details      = {  
  id      = "my_iam_user_vino"  
  ~ name   = "my_iam_user_vino" -> "my_iam_user_vinoth"  
  ~ tags   = null -> {}
```

```
    # (6 unchanged attributes hidden)
  }
```

#command

terraform apply iam.tfplan

#output

aws_iam_user.my_iam_user: Modifying... [id=my_iam_user_vino]

aws_iam_user.my_iam_user: Modifications complete after 1s

[id=my_iam_user_vinoth]

Apply complete! Resources: 0 added, 1 changed, 0 destroyed.

Instead of doing the entire file check and update, just to update the user alone we could do below command.

#command

terraform apply -target=aws_iam_user.my_iam_user

Terraform intelligently understands that S3 bucket once created cannot be renamed so it deletes it and creates a new one with a given new name, whereas iam username can be renamed so it does not delete the user, it renames it.

Terraform States files:

Whenever **terraform apply** is performed,

1. **terraform.tfstate** - state after successful execution of command will be saved in **terraform.tfstate** file
2. **terraform.tfstate.backup** - state before execution of command will be saved in **terraform.tfstate.backup** file

If there is no known state i.e.(**terraform.tfstate** file), it cannot compare known state with the actual state.

If we try to run a plan, it will try to **create a user and s3 bucket as per main.tf file.**

Let's understand the reason...

Terraform works mainly on the **terraform name.**

Terraform matches the **terraform name** from **main.tf** resources with the name from the **terraform.tfstate** file. And it takes the **unique_id** from the **tfstate** and tries to match it with **AWS cloud.**

If **terraform.tfstate** is not present terraform will have no idea, so it will try to create new resources.

So known state(**terraform.tfstate**) is very informant.

If in a project if multiple people are working, we must share them the known state.

Terraform state is an unencrypted file containing sensitive information committing it in git or version control repo is un advisable.

We can use something called a **remote backend** in cloud like **S3**

#excluding the files for committing in git

```
# terraform excludes
*.tfstate
*.tfstate.backup
.terraform
```

Creating all the resources in a single **main.tf** file will make it complex.

The file can be any name not just **main.tf** it can be any name if **tf** as extension.

All **tf** extension files are concatenated and executed at a single time.

Destroy All the resources created with Terraform.

1. Resources created is **IAM user** and **S3 bucket**.

#command to destroy all the resources
terraform destroy

vinoth\$ terraform destroy

Output:

aws_iam_user.my_iam_user: Refreshing state... [id=my_iam_user_vinoth]
aws_s3_bucket.my_s3_bucket: Refreshing state... [id=my-s3-bucket-vinothkumar-7]

terraform destroy refreshes the state, looks the current state, identifies the resources, and deletes them.

Do you really want to destroy all resources?

Terraform will destroy all your managed infrastructure, as shown above.
There is no undo. Only 'yes' will be accepted to confirm.

Enter a value: yes

aws_iam_user.my_iam_user: Destroying... [id=my_iam_user_vinoth]
aws_s3_bucket.my_s3_bucket: Destroying... [id=my-s3-bucket-vinothkumar-7]

aws_iam_user.my_iam_user: Destruction complete after 1s
aws_s3_bucket.my_s3_bucket: Destruction complete after 1s

Destroy complete! Resources: 2 destroyed.

Terraform to Create multiple IAM Users using Terraform:

mentioning the count and the index of count. This is supported by HCL.

```
resource "aws_iam_user" "my_iam_user" {  
    count = 2  
    name = "my_iam_user_vinoth_${count.index}"  
}
```

#initializing the new project to download necessary providers.
terraform init

#save the plan of iam policy
terraform plan -out iam.tfplan

#execute the plan
terraform apply iam.tfplan

#output
aws_iam_user.my_iam_user[1]: Creating...
aws_iam_user.my_iam_user[0]: Creating...
aws_iam_user.my_iam_user[0]: Creation complete after 1s
[id=my_iam_user_vinoth_0]
aws_iam_user.my_iam_user[1]: Creation complete after 1s
[id=my_iam_user_vinoth_1]

#command to validate the terraform file
terraform validate

#command to format the terraform file
terraform fmt

#command to show the current state of resources
terraform show

Terraform helps more in Error validation. If there are scripts with errors it handles them accordingly.

Variables in Terraform:

Variables are important because they make the terraform scripts dynamic.

```
variable "iam_user_name_prefix" {  
  type = string  
  default = "my_iam_user"  
}
```

```
resource "aws_iam_user" "my_iam_user" {  
  count = 3  
  name = "${var.iam_user_vinoth_prefix}_${count.index}"  
}
```

#command
terraform apply -refresh=false

#output
Enter a value: yes

aws_iam_user.my_iam_user[2]: Creating...
aws_iam_user.my_iam_user[2]: Creation complete after 2s
[id=my_iam_user_vinoth_2]

Apply complete! Resources: 1 added, 0 changed, 0 destroyed.

1. Variable type by default its **any**. Unless it's explicitly mentioned.
#any, number, bool, list, map, set, object, tuple.

2. If default value is not given for any variable, terraform will ask for a value during **validate**.

3. Variable value can be **overridden** from the **default value** to **env variable** value.

export

TF_VAR_iam_user_name_prefix=FROM_ENV_VARIABLE_IAM_PREFIX

Note: **terraform.tfvars** name is default for variable file if needed it can be overridden by **terraform apply -var-file="some-name.tfvars"**

4. Another way of overriding the variable value is by creating a file **terraform.tfvars** and provide the value from the file as below.

iam_user_name_prefix="VALUE_FROM_TERRAFORM_VARS"

output:

name = "my_iam_user_vinoth_0" ->
"VALUE_FROM_TERRAFORM_VARS_0"

Plan: 0 to add, 1 to change, 0 to destroy.

5. Another way of overriding the variable value is by giving value from **CMD line**

terraform plan -refresh=false -

var="iam_user_name_prefix=VALUE_FROM_COMMAND_LINE"

Priority for variable value:

Value from CMD > terraform.tfvars > FROM_ENV_VARIABLES > Default value.

List and Sets Variables:

1. Creating list of usernames

```
variable "names" {  
  default = ["ranga","tom","jerry"]  
}
```

```
resource "aws_iam_user" "my_iam_user" {  
  count = length(var.names)  
  name = var.names[count.index]  
}
```

#output

Plan: 3 to add, 0 to change, 0 to destroy.

Do you want to perform these actions?

Enter a value: yes

```
aws_iam_user.my_iam_user[1]: Creating...  
aws_iam_user.my_iam_user[2]: Creating...  
aws_iam_user.my_iam_user[0]: Creating...  
aws_iam_user.my_iam_user[0]: Creation complete after 1s [id=ranga]  
aws_iam_user.my_iam_user[2]: Creation complete after 1s [id=jerry]  
aws_iam_user.my_iam_user[1]: Creation complete after 1s [id=tom]
```

Apply complete! Resources: 3 added, 0 changed, 0 destroyed.

terraform console

There is no function named "count".

```
> length(var.names)
3
> reverse(var.names)
[
  "jane",
  "tom",
  "ranga",
]
> distinct(var.names)
[
  "ranga",
  "tom",
  "jane",
]
> toset(var.names)
[
  "jane",
  "ranga",
  "tom",
]
> █

> concat(var.names, ["new_value"])
[
  "ranga",
  "tom",
  "jane",
  "new_value",
]
> contains(var.names, "ravi")
false
> contains(var.names, "ranga")
true
> sort(var.names)
[
  "jane",
  "ranga",
  "tom",
]
_
```

```
variable "names" {
  default = ["ranga", "tom", "jerry"]
}
```

```
resource "aws_iam_user" "my_iam_user" {
  count = length(var.names)
  name  = var.names[count.index]
}
```

Whenever we change the order in list in the list of names we added above,

```
variable "names" {  
  default = ["sata", "ranga", "tom", "jerry"]  
}
```

Now as we added Sata in the beginning of the list

When we apply the changes

#output

Plan: 1 to add, 3 to change, 0 to destroy.

Because, we are trying to change the position in the list, all the users have to be shifted.

Whenever we add the variable with **list or count**, **terraform** stores the variables as **list of elements**. Each of them are **indexed** as list. List index compares with each element trying to change all.

We can overcome this by using foreach

```
variable "names" {  
  default = ["sata", "ranga", "tom", "jerry"]  
}
```

```
resource "aws_iam_user" "my_iam_user" {  
  #count = length(var.names)  
  #name = var.names[count.index]  
  for_each = toset(var.names)  
  name = each.value  
}
```

Variables can be iterated using **foreach** only if it's **unique**. So we are **converting it into set** and iterating the values.

Now even if we add new variables, it won't be a problem.

Index will be the variables names.

NOTE:

1. Deletion and updating will be based on the **index in case of using count**
2. Deletion and updating will be based on the **variable name in case of using foreach.**

Variable Names with Sets:variable "users" {

```
    default = {
        ravs:"Netherlands",
        tom:"US",
        jane:"Inidia"
    }
}
```

```
resource "aws_iam_user" "my_iam_users" {
    for_each = var.users
    name = each.key
    tags = {
        country: each.value
    }
}
```

```
MacBook-Pro:04-terraform-maps vinoth$ terraform console
> var.names
{
  "jane" = "Inidia"
  "ravs" = "Netherlands"
  "tom" = "US"
}
> var.names.jane
"Inidia"
> keys(var.names)
[
  "jane",
  "ravs",
  "tom",
]
> values(var.names)
[
  "Inidia",
  "Netherlands",
  "US",
]
```

Output:

Enter a value: yes

```
aws_iam_user.my_iam_users["jane"]: Creating...
aws_iam_user.my_iam_users["tom"]: Creating...
aws_iam_user.my_iam_users["ravs"]: Creating...
aws_iam_user.my_iam_users["tom"]: Creation complete after 1s [id=tom]
aws_iam_user.my_iam_users["jane"]: Creation complete after 1s [id=jane]
aws_iam_user.my_iam_users["ravs"]: Creation complete after 1s [id=ravs]
```

Maps of maps:

Adding country

```
variable "users" {
  default = {
    ravs: { country: "Netherlands", department: "BDA"},
    tom: { country: "US", department: "ADS"},
    jane: { country: "India", department: "Ds"}
  }
}
```

```
resource "aws_iam_user" "my_iam_users" {
  for_each = var.users
  name = each.key
  tags = {
    country: each.value.country
    department: each.value.department
  }
}
```

Adding another variable department

```
variable "users" {
  default = {
    ravs: { country: "Netherlands", department: "BDA"},
    tom: { country: "US", department: "BDSA"},
    jane: { country: "India", department: "BSDVDA"}
  }
}
```

```

    }
}

resource "aws_iam_user" "my_iam_users" {
  for_each = var.users
  name = each.key
  tags = {
    country: each.value.country
    department: each.value.department
  }
}

```

EC2 in Amazon Console (EC2 is virtual server in AWS Cloud)

Service -> EC2 -> Launch Instance -> choose the below configurations.

- 1. Choose region**
- 2. Choose AMI**
- 3. Choose VPC**
- 4. Choose Subnet**
- 5. Choose Storage**
- 6. Choose security group**
- 7. Tags in any**
- 8. Launch Instance**

Regions:

By having multiple regions we can provide high availability and low latency for users. .

Availability Zone:Each region have multiple availability zone. AZ's are within the regions but are physically separated data center. (For availability)

VPC:

Virtual Private Cloud is firewall for cloud. It has subnets.

Subnet:

Subnet can be **private** or **public**.

Security Group:Way to control traffic to EC2 instance. Additional traffic control other than Subnet for specific Ec2 instances.

Ingress - From where the traffic should be allowed from.

Egress - What can be done from the HTTP server.

A default security group given will have default egress for everything(to any system on the ip)

VM - Virtual Server

EC2 - Elastic Compute Cloud

US EAST (N. Virginia) us-east-1

ami-0a699202e5027c10d

t2.micro

vpc-0e1388f6633187dae

Steps:

1. Choose a region where we want to create the EC2 instance in.
US EAST (N. Virginia) us-east-1
2. Service -> EC2 -> Launch Instance -> choose the below configurations.
3. We have to choose Amazon Machine Image(AMI) (OS and the software to be present in the server.)-> ami-0a699202e5027c10d
4. Choose the instance type (type, cpu, Memory, Storage, network performance, I-support) -> t2.micro
5. Choose the **network** and **Subnet**:I. Whenever we create a resource in data center they are already protected from a **physical firewall**. For cloud we have to create a **Virtual Private Cloud(VPC)**.II. In **VPC** we can create multiple **Subnets**.
- If subnet is private, only the resources within the VPC will be able to talk to the subnet and not from outside
(**Databases** that shouldn't be allowed to talk to outside the from VPC we can put into **private subnets**)
- If subject is public, for the resources can talk with VPC and send requests (**Webservers**, can be put into public subnets)
6. For **each region** a **VPC** is given by **default** - vpc-0e1388f6633187dae

7. Choose **storage**
8. Create **security group**. For SSH, HTTP, HTTPS and give custom ip addresses.

Creating EC2 instance Using Terraform

Service -> EC2 -> Network & Security -> Security Groups

1. Creating a security group.

Egress has to be explicitly mentioned in terraform. It is not configured by default.

1. If there is a change in name of the security group, it will delete the existing resource and create a new resource.

```
provider "aws" {  
  region = "us-east-1"  
}
```

Http server -> 80 TCP, SSH 22 TCP, CIDR(Used to specify range of ip addresses.) ["0.0.0.0/0"]

```
resource "aws_security_group" "http_server_sg" {  
  name = "http_server_sg"  
  vpc_id = "vpc-0e1388f6633187dae"
```

Ingress for HTTP

```
  ingress {  
    from_port = 80  
    to_port   = 80  
    protocol  = "tcp"  
    cidr_blocks = ["0.0.0.0/0"]  
  }
```

Ingress for SSH

```
  ingress {  
    from_port = 22
```

```

    to_port    = 22
    protocol   = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  # Egress for all traffic
  egress {
    from_port = 0
    to_port   = 0
    protocol  = "-1" # All protocols
    cidr_blocks = ["0.0.0.0/0"]
  }
  #tags are used to tie the resource to a specific environment which helps in
  identification
  tags = {
    name = "http_server_sg"
  }
}

```

Output:

Plan: 1 to add, 0 to change, 0 to destroy.

```
terraform apply "ec2plan.tfplan"
```

MacBook-Pro:05-ec2-instances vinoth\$ terraform apply ec2plan.tfplan

aws_security_group.http_server_sg: Creating...

aws_security_group.http_server_sg: Creation complete after 4s [id=sg-0343ad64e36d92184]

Apply complete! Resources: 1 added, 0 changed, 0 destroyed.

2. Creating a Key Pair

Service -> EC2 -> Network & Security -> Key Pair

Key pair is required to run command from SSH to EC2 instances.

1. Create key pair
2. Name : default-ec2
3. Key pair type: RSA

4. Key file format: .pem (For use with openSSH)
5. Created.

#commands to save the file in system to protect it

```
Chmod 777 default-ec2.pem
```

```
chmod 400 default-ec2.pem
```

```
mkdir ~/aws
```

```
MacBook-Pro:Downloads vinoth$ mkdir ~/aws/aws_keys
```

```
MacBook-Pro:Downloads vinoth$ mv default-ec2.pem ~/aws/aws_keys
```

Creating a Ec2 Instance(http Server) in Terraform:

```
resource "aws_instance" "http_server" {  
  ami = "ami-0a699202e5027c10d"  
  key_name = "default-ec2"  
  instance_type = "t2.micro"  
  vpc_security_group_ids = [aws_security_group.http_server_sg.id]  
  subnet_id = "subnet-0f95b394e50f22512"  
}
```

#output

```
MacBook-Pro:05-ec2-instances vinoth$ terraform apply ec2plan.tfplan
```

```
aws_instance.http_server: Creating...
```

```
aws_instance.http_server: Still creating... [10s elapsed]
```

```
aws_instance.http_server: Still creating... [20s elapsed]
```

```
aws_instance.http_server: Still creating... [30s elapsed]
```

```
aws_instance.http_server: Creation complete after 35s [id=i-  
0fef5f1d42fcb537]
```

Connecting to the HTTP server using already created Key pair

```

variable "aws_key_pair" {
  default = "~/aws/aws_keys/default-ec2.pem"
}

resource "aws_instance" "http_server" {
  ami           = "ami-0a699202e5027c10d"
  key_name      = "default-ec2"
  instance_type = "t2.micro"
  vpc_security_group_ids = [aws_security_group.http_server_sg.id]
  subnet_id     = "subnet-0f95b394e50f22512"

  connection {
    type      = "ssh"
    host      = self.public_ip
    user      = "ec2-user"
    private_key = file(var.aws_key_pair)
  }

  provisioner "remote-exec" {
    inline = [
      #install httpd, start, copy a file
      "sudo yum install httpd -y",
      "sudo service httpd start",
      "echo Virtual server is at ${self.public_dns} | sudo tee
/var/www/html/index.html"
    ]
  }
}

```

By applying the changes done above, will not make any changes in EC2.

#output terraform apply

aws_security_group.http_server_sg: Refreshing state... [id=sg-0343ad64e36d92184]
aws_instance.http_server: Refreshing state... [id=i-0fef5f1d42fcb537]
No changes. Your infrastructure matches the configuration.
Apply complete! Resources: 0 added, 0 changed, 0 destroyed.

After creating EC2 instances, any changes done after creation is **immutable**.

To make the changes effective we have to **destroy** the **ec2** and create it again.

#it will destroy the server first and the security group secondly because server depends on security group.

terraform destroy

#output

Plan: 0 to add, 0 to change, 2 to destroy.
Destroy complete! Resources: 2 destroyed.

terraform apply

Apply complete! Resources: 2 added, 0 changed, 0 destroyed.

Immutability and Why Immutable Server?

While provisioning servers using IAC, If new server is needed with a current state. It is recommended to create a new server with version as v2. Once v2 is up and running we can remove the old v1 server.

While using IAC, immutable servers are recommended and used.

Remove hardcoded variable values in Terraform

aws_default_vpc - is managed by aws.

```
resource "aws_default_vpc" "default" {  
  
}
```

```
terraform apply -target=aws_default_vpc.default  
#output
```

Plan: 1 to add, 0 to change, 0 to destroy.

Apply complete! Resources: 1 added, 0 changed, 0 destroyed.

#default value harcoding can be avoided by this. Now the vpc_id is dynamic

```
resource "aws_security_group" "http_server_sg" {  
  name = "http_server_sg"  
  //vpc_id = "vpc-0e1388f6633187dae"  
  vpc_id = aws_default_vpc.default.id
```

#removing the hard coding value of subnets using data provider

```
data "aws_subnet" "default_subnets" {  
  vpc_id = aws_default_vpc.default.id  
}
```

```
#subnet_id = data.aws_subnet.default_subnets.id
```

#removing the hard coding value of AMI

```
data "aws_ami" "aws_linux_2_latest" {  
  most_recent = true  
  owners = ["amazon"]  
  filter {  
    name = "name"  
    values = ["amzn2-ami-hvm-*"]  
  }  
}
```

```
resource "aws_instance" "http_server" {  
  ami = "ami-0a699202e5027c10d"  
  #ami = data.aws_ami.aws_linux_2_latest
```

```
terraform apply -target=data.aws_ami.aws_linux_2_latest
```

Terraform Graph (Dye Graph)

It is the resource graph.

Remote Backend for Storing the Tf States - S3

1. Store all the state of all projects in S3 bucket.
2. Locking - To avoid Concurrency between users trying to use state with help of DynamoDB for locking and isolation.
 - I. Acquire a lock.
 - II. Update the resources.
 - III. Release the lock.
3. Encryption.

Partition the resources into two folders:

07-backend-state/backend-state -> Where all the bucket configuration gonna take place

07-backend-state/users - > All user projects and other details.

07-backend-state/backend-state/main.tf:

```
provider "aws" {  
  region = "us-east-1"  
}
```

```

resource "aws_s3_bucket" "getajob_backend_state" {
  bucket = "application-name-backend-state-vinoth"
  lifecycle {
    prevent_destroy = true
  }
  versioning {
    enabled = true
  }
  server_side_encryption_configuration {
    rule {
      apply_server_side_encryption_by_default {
        sse_algorithm = "AES256"
      }
    }
  }
}

```

```

resource "aws_dynamodb_table" "enterprise_backend_lock" {
  name      = "dev_application_locks"
  billing_mode = "PAY_PER_REQUEST"
  hash_key  = "LockID"

  attribute {
    name = "LockID"
    type = "S"
  }
}

```

Terraform apply
#output

Apply complete! Resources: 2 added, 0 changed, 0 destroyed.

07-backend-state/users/main.tf:

```

variable "application_name" {
  default = "01-backend-state"
}

```

```

variable "project_name" {
  default = "users"
}

variable "environment" {
  default = "dev"
}

terraform {
  backend "s3" {
    bucket = "application-name-backend-state-vinoth"
    # key    = "${var.application_name}-${var.project_name}-${var.environment}"
    key      = "application_name-project_name-environment"
    region   = "us-east-1"
    dynamodb_table = "dev_application_locks"
    encrypt   = true
  }
}

provider "aws" {
  region = "us-east-1"
}

resource "aws_iam_user" "my_iam_user" {
  name = "my_iam_user_vinoth"
}

```

Now the keys will be saved in s3.

Creating multiple environments using Terraform Workspace:

All details for an environment it can be stored in workspace.

To switch between workspace :

```
terraform workspace select default
```

terraform workspace

#Output

-default

#creating a new workspace

terraform workspace new prod-env

This is create the environment in s3 bucket.

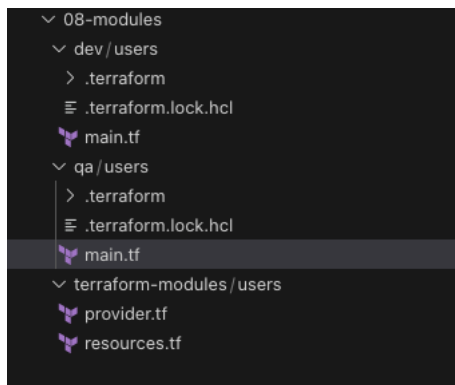
#to list all the workspace

terraform workspace list

#to show current workspace environment

terraform workspace show

Modules in Terraform:



```
terraform > 08-modules > dev > users > main.tf
1  module "user_module" {
2      source = "../../terraform-modules/users"
3      environment = "dev"
4  }
```

```
terraform > 08-modules > qa > users > main.tf
1  module "user_module" {
2      source = "../../terraform-modules/users"
3      environment = "qa"
4  }
```

```

terraform > 08-modules > terraform-modules > users > resources.tf
1  variable "environment" {
2      default = "default"
3  }
4  resource "aws_iam_user" "my_iam_user"{
5      name = "my_iam_user_abc_${var.environment}"
6  }
7  }

```

Variables are two types:

1. Global
2. Local - variables cannot be overridden

```

terraform > 08-modules > terraform-modules > users > resources.tf
1  #global variable
2  variable "environment" {
3      default = "default"
4  }
5  #local variable
6  locals {
7      iam_user_extension = "my_iam_user_abc"
8  }
9  resource "aws_iam_user" "my_iam_user"{
10     name = "${local.iam_user_extension}_${var.environment}"
11 }
12

```